

Provenance

Every design decision in the tool, where to see it, why it was made, and the source. Current as of v0.4. One row per decision; anything needing a paragraph is in design-rationale.md.

Rows marked **ours** are not from any paper. They are listed because a decision with no source is worth knowing about too.

Sources

SHORT NAME	CITATION
Driscoll	Driscoll, Chen, Shi, Vucharatavintara, Yao, Jin. "Understanding Parents' Desires in Moderating Children's Interactions with GenAI Chatbots through LLM-Generated Probes." CHI '26. 10.1145/3772318.3791622
iRULER	Bai, Cheong, Muller, Lim. "iRULER: Intelligible Rubric-Based User-Defined LLM Evaluation for Revision." CHI '26. 10.1145/3772318.3790539
PolicyPad	Feng, Kuo, Chen, Cheong, Holstein, Zhang. "PolicyPad: Collaborative Prototyping of LLM Policies." CHI '26. 10.1145/3772318.3791689
Botender	Kuo, Liu, Chen, Seering, Zhang, Zhu, Holstein. "Botender: Supporting Communities in Collaboratively Designing AI Agents through Case-Based Provocations." CHI '26. 10.1145/3772318.3790500
ICAI	Findeis, Kaufmann, Hullermeier, Albanie, Mullins. "Inverse Constitutional AI: Compressing Preferences into Principles." ICLR 2025
Ang	Ang, Gollapalli, Ng. EACL 2023, Table 1. Reproduces the Paul and Elder question types
Robinson	Robinson. "Probing in qualitative research interviews." Qualitative Research in Psychology 20, 3 (2023). Driscoll's probing method

What the participant sees

What the participant sees			
DECISION	WHERE YOU SEE IT	WHY	SOURCE
1 Two reflective questions are asked at setup, before any case exists	The page after sign-up, under the child's age	Every other question in the tool hangs off a situation we generated, so the first thing a participant considers is our material. These come first and are about them and their child, so what they bring to the cases is their own. Min: "it might be even before they see the specific case... maybe add a few more questions when you ask about name and tell me about your kid"	<i>Min's meeting, 22 Aug</i>
2 The setup page asks only the child's age and, optionally, their name	Same page	The age is already collected at sign-up, so asking who the agent is for in free text asked the same thing twice. The name, when given, goes into the generated situations so they read "Maya asks" rather than "the child asks"	<i>ours, from Pranav's review</i>
3 Questions carry no type label and no card around them	Every question in the tool	Each sat in a bordered card with a pill naming its Socratic type, a provenance icon and a red "needed" tag: four pieces of chrome around one sentence, and the type names our taxonomy rather than the participant's task. The type is still recorded in the data	<i>ours, from Pranav's review</i>
4 Validation counts questions rather than naming them	Under the save and next buttons	The message used to list the Socratic type names, which pointed at nothing on screen once the labels were removed	<i>ours</i>
5 Topics come from eight named parent concerns, not from us	The theme menu, eight cards	Parents were interviewed about real child-chatbot transcripts and their worries were coded into eight themes. Using them means the tool covers what parents actually raised rather than what we assumed	<i>Driscoll Table 3</i>
6 Each theme shows a plain description and how many of 24 parents raised it	Under each theme name	The paper's labels ("Wrong Approach to Delivery") mean nothing to someone who has not read it, and the count is what makes the menu read as findings rather than our opinion	<i>Driscoll Table 3</i>
7 One rule per theme, written across three cases	The workspace: three cases, one rule box	A rule written for a single case is a local reaction. Writing one rule that must hold across three makes it a general principle with concrete examples under it	<i>ours, from Min's review</i>

What the participant sees				
	DECISION	WHERE YOU SEE IT	WHY	SOURCE
8	Cases are generated per theme, three at a time	Appear when you open a theme	Generating against one named theme lets the category be set in code rather than inferred, so a case cannot land in a bucket the menu cannot reach	<i>Botender A.2.4 (set-level generation)</i>
9	A case is dropped if the person's stated wishes already settle it	Not visible; a filter in the prompt	A case whose answer is already determined teaches nothing. This is the mechanism behind surfacing what someone had not considered	<i>Botender A.2.1</i>
10	No trivial cases, and each must be self-evident from its own text	Case quality, throughout	A case that turns on tone alone, or that needs explaining, does not provoke a real decision	<i>Botender A.2.1-A.2.3</i>
11	The agent can only allow, soften, add context, decline, or involve the parent	Case generation prompt	Cases that need abilities the agent does not have cannot be answered	<i>Botender A.2.5</i>
12	Letting content through is a legitimate answer	Baseline replies and options	Without this the tool only ever elicits restriction	<i>Botender A.1.2</i>
13	The case card shows the title and situation, nothing else	The workspace	It used to also show what was at stake, the tradeoffs, and how it plays out, which handed the participant the answer before they wrote one	<i>Min's review</i>
14	What is at stake and the competing values become questions, not statements	The "Before you answer" questions	Same content, opposite effect: as prose it tells them what to think, as a question it makes them think	<i>Min's review</i>
15	Every case appears as an actual conversation	Case tabs, chat bubbles	A description of a moment and the moment itself are different things to react to. Two layers, a synopsis plus the exchange	<i>PolicyPad §5.2.2, Fig. 5D</i>
16	The "before" reply is a real product default, not a bare model	The agent's turn in every case	It applies the safeguards a mainstream child mode ships with: age-appropriate wording, pointing to a trusted adult, no graphic detail, a crisis resource on self-harm. Reconstructed from published behaviour; we do not have anyone's real prompt	<i>ours</i>

What the participant sees				
	DECISION	WHERE YOU SEE IT	WHY	SOURCE
17	That reply has no length limit	Same	It was capped at two to four sentences, which made the before-state read as unrepresentative next to a real product. Length now varies with the question	<i>Min's review</i>
18	Questions before writing use five question types designed to provoke examination	"Before you answer"	These types exist to make someone examine a thought rather than defend it, and they have no correct answer	<i>Ang Table 1 (Paul and Elder)</i>
19	A pre-write question may never attribute a belief to the participant	The question prompt's HARD RULE	One asked why the parent assumed something the AI's own case had introduced. Attributing a position to someone who has not taken one corrupts what they write next	<i>Min's review</i>
20	Each case carries one counterfactual "what if" question	"As you write, consider"	An interviewer does not take a first answer; they change one detail to see whether the reasoning holds. Ours is fixed at generation time and cannot adapt the way a live interviewer would	<i>Driscoll §3.5, method from Robinson</i>
21	Questions are required, not optional	Save is blocked until answered	Their purpose is to make someone think before writing, and they feed the rubric check, so a skipped one silently weakens it	<i>Min's review</i>
22	The rule is scored on four weighted criteria with four levels each	"How your rule scores"	A rule can be exactly right about what someone wants and still be too vague for an agent to follow. The rubric checks only that narrower thing	<i>iRULER §4.1.1, App. B</i>
23	The four criteria's content comes from what parents said mattered	The criteria names	Specific action, scope, developmental fit, escalation are a synthesis of their coded moderation themes, not a one-to-one mapping	<i>Driscoll §4.2, §5.3</i>
24	Every criterion shows why it scored, and what the next level needs	The expanders under the score	Feedback that says only "level 2" is not actionable. The why-not-higher was their most used feature	<i>iRULER A.1.1, Table 4</i>
25	The rubric never sees the reflection answers as scorable text	Not visible; a fence in the prompt	With reflection added, a draft naming nobody scored level 4 on "Fits the audience" in 2 of 3 runs, because the model credited intent the rule did not contain	<i>ours, found in testing</i>

What the participant sees				
	DECISION	WHERE YOU SEE IT	WHY	SOURCE
26	No AI rewrites the rule	There is no suggest-a-revision button	It existed and was removed. In the study it came from, participants accepted every suggestion offered, and it is one more generated component needing accuracy validation	<i>iRULER §7.3.1 acceptance rate; removal was Min's call</i>
27	Instead, the participant's own words are quoted back	"Put this in your rule"	If they said something while thinking and their rule does not mention it, the gap is theirs to close in their own language	<i>Min's review</i>
28	The rule can be tested against the conversation at any time	"Try it on a case"	Writing a rule and immediately seeing what it does is the fastest way to find out whether it says what was meant	<i>PolicyPad §5.2.5, Fig. 8</i>
29	The rule is delivered to the model as a system prompt over the defaults	Not visible	The parent's instructions sit on top of the product defaults and win where they disagree	<i>PolicyPad §5.3</i>
30	Two rule versions can be asked the same question side by side	"Compare versions"	Seeing one difference at a time is what makes a change legible	<i>Min's review</i>

Testing what was written

Testing what was written				
	DECISION	WHERE YOU SEE IT	WHY	SOURCE
1	Three rounds run per theme, and none at the end	After saving a rule	The theme rounds test whether that rule predicts the person's choices. Min cut the all-themes round: after the last theme the session goes straight to export	<i>Min's review (Sep 1)</i>
2	Each comparison changes exactly one thing between the two replies	The "varies" tag on each moment	If two things change, a choice says nothing specific. The four dimensions are the child's age, severity, clarity of intent, and where the risk comes from	<i>Driscoll §7 and Table 14</i>

Testing what was written				
	DECISION	WHERE YOU SEE IT	WHY	SOURCE
3	A pick is committed with a reason before anything is revealed	Round 2 and 3	Seeing the model's answer first would contaminate the reason, and the agreement score would stop measuring anything	<i>ours</i>
4	The buttons come before the reason box, and nothing auto-advances	All rounds	The reason box used to sit above the buttons, asking people to justify a choice they had not made, and one click skipped past it entirely	<i>Min's review</i>
5	In the theme rounds the model follows only that theme's rule	Round 2's reveal	A disagreement then points at a gap in that rule, rather than at an aggregate the person never wrote	<i>ours</i>
6	Round 2 hands the rule box back rather than offering a rewrite	Under the reveal	Same reason the AI revision was removed. Round 3 then scores the edited rule, so the loop measures whether their own revision closed the gap	<i>ours</i>
7	Round 3 is frozen and scored	The last round	Agreement means nothing if the thing being measured changes while measuring it	<i>ours</i>
8	The workspace is the split layout, staged	Every theme	Min picked it from five prototypes: the case stays pinned left while the work steps down the right	<i>Min's review (Sep 1)</i>
9	The rule box sits directly below the questions	Consider + write	Min's fix for "they cannot see their answer": the reflections stay on screen while the first draft is written	<i>Min's review (Sep 1)</i>
10	Testing is the last stage, with versions and a final pick	The Test stage	People iterate by chatting against the rule; the final version is chosen explicitly because it is not always the last one	<i>Min's review (Sep 1)</i>
11	Every rule box carries a version dropdown	Score + revise, Sharpen, Test	It was unclear which version was in the box or produced a reply; Original is the no-rule baseline and v1 is the first draft	<i>user direction (Sep 2)</i>
12	Each version keeps its own persistent chat thread	The Test stage's compare columns	Switching versions switches threads and nothing is lost, so a reply always traces to the version that produced it	<i>user direction (Sep 2)</i>

Testing what was written				
	DECISION	WHERE YOU SEE IT	WHY	SOURCE
13	Three side-by-side compare columns replace the single chat	The Test stage	Pick a version per column and ask all of them the same things; the case's own question seeds an empty thread	<i>user direction (Sep 2)</i>
14	A round-2 edit is re-recorded immediately	The theme rounds	Otherwise the export and later rounds keep the pre-edit rule while the person believes they revised it	<i>ours</i>

Study mechanics

Study mechanics				
	DECISION	WHERE YOU SEE IT	WHY	SOURCE
1	No nested scrolling anywhere in the workspace	Every layout	Panes with a fixed height smaller than their own contents put a second scrollbar inside the page, so on one layout the answer boxes sat below the fold of a pane with no sign they were there. The only pane that still scrolls is the test transcript, which grows without bound	<i>ours, from Pranav's review</i>
2	Every generated artifact carries an info icon	The small "i" beside each	A participant can see what produced anything on screen and where the method came from. Also how we keep ourselves honest about what is borrowed	<i>ours</i>
3	Session state is snapshotted after every interaction	Not visible; the resume code	Someone can leave and come back, and a crash does not lose a session	<i>ours</i>
4	Sessions saved before the theme restructure reset to the theme menu	Only if resuming an old code	Per-case answers cannot honestly be reshaped into per-theme ones, so the old work is cleared with an explanation rather than half-restored	<i>ours</i>
5	Four layout prototypes, switchable	A control visible only in test sessions	Min asked to see the options rather than have them described. Not a participant-facing feature	<i>Min's request</i>

Study mechanics				
	DECISION	WHERE YOU SEE IT	WHY	SOURCE
6	Study arms exist but randomisation is off	Not visible	Handing a random visitor a version with no rule-writing screen is not a default we want. <code>STUDY_ARMS=on</code> enables it	<i>ours</i>

Convergent, not borrowed

One row for honesty in the other direction. Putting a published taxonomy into a generator prompt was already in the tool before Botender was read; Botender A.2.3 does the same thing with a different taxonomy. What we did take from that section is the convention of citing the technique and the embedded taxonomy separately.

Known gaps

- **Which dimension a comparison varies is a model judgment**, not a selector targeting expected disagreement or cycling for coverage. It is recorded in each row so coverage can be analysed after the fact. Min called dimension choice "the key".
- **Nothing here has been validated with human coders**. Botender ran a 90-participant study on its case generation; iRULER had two experts rate 96 essays. We have run neither.
- **The child-safe baseline is a reconstruction** from published behaviour, not any vendor's actual instructions.