

# What we took, and from where

The papers behind the Scenario Elicitor, and concretely what each one gave us. The full decision-by-decision record is in provenance.md; this is the short version, one entry per paper.

## Driscoll et al. — the topics, the criteria, and what the comparisons vary

---

Driscoll, Chen, Shi, Vucharatavintara, Yao, Jin. "Understanding Parents' Desires in Moderating Children's Interactions with GenAI Chatbots through LLM-Generated Probes." CHI '26. 10.1145/3772318.3791622

They showed 24 parents real conversations between children and AI chatbots and coded what worried them. We use their results in three places.

**The eight topics a participant chooses from** are their Table 3 verbatim, including the count of how many parents raised each. Five are about what the AI said back, three about what the child asked. Nothing in the topic list is our guess about what parents care about.

**The four things the rubric scores** are a synthesis of their coded moderation themes: naming a specific action, setting scope, fitting the child's stage, and saying when to escalate. This is interpretation rather than extraction; no table in the paper contains these four.

**The four dimensions a comparison varies** are theirs. Three are named in their future work (clarity of intent, severity, developmental cues) and the fourth, where the risk originates, is their system-risk against misuse-risk split.

## iRULER — the shape of the rubric

---

Bai, Cheong, Muller, Lim. "iRULER: Intelligible Rubric-Based User-Defined LLM Evaluation for Revision." CHI '26. 10.1145/3772318.3790539

**The rubric's structure:** named criteria, four ordinal levels each described as observable evidence, and a why-not-higher explanation attached to every score. Their instrument is three unweighted booleans over a whole document; ours is four weighted criteria over one rule.

**One thing we took as a warning rather than a feature.** They offer an AI revision targeting the weakest criterion, and their participants accepted 100% of the suggestions offered. We built it, then removed it, because a channel that reliably turns the model's language into the person's stated preference.

## Ang et al. — the reflective questions

---

Ang, Gollapalli, Ng. EACL 2023, Table 1. Reproduces the Paul and Elder taxonomy.

**The five question types** every reflective question in the tool is generated from: clarification, probing assumptions, probing reasons and evidence, probing implications, and probing alternative viewpoints. These questions have no correct answer; they exist to make someone examine a thought rather than

defend it. The type is recorded with each answer and never shown to the participant, because it names our taxonomy rather than their task.

## PolicyPad — how a rule is shown, delivered, and tested

---

Feng, Kuo, Chen, Cheong, Holstein, Zhang. "PolicyPad: Collaborative Prototyping of LLM Policies." CHI '26. 10.1145/3772318.3791689

**Two-layer cases:** a short synopsis plus the actual conversation, rather than a description of a situation on its own.

**The rule reaches the model as a system prompt** with scaffolding, and the precedence line saying the person's instructions win is modelled on their open-source scaffold.

**Write, test, revise, test again.** Both of their testing paths: regenerate the reply, or continue the conversation. Ours versions one rule for one topic where theirs versions a whole policy document.

## Botender — what makes a case worth showing

---

Kuo, Liu, Chen, Seering, Zhang, Zhu, Holstein. "Botender: Supporting Communities in Collaboratively Designing AI Agents through Case-Based Provocations." CHI '26. 10.1145/3772318.3790500

**The quality bar every generated case has to clear.** Drop a case whose answer is already settled by what the person has told us, since it teaches nothing. No cases that turn on tone alone. Each must be understandable from its own text. And the agent can only do what it can actually do: allow, soften, add context, decline, or bring the parent in.

**Non-action counts as an answer**, so the tool does not only ever elicit restriction.

## Inverse Constitutional AI — why we measure agreement at all

---

Findeis, Kaufmann, Hullermeier, Albanie, Mullins. "Inverse Constitutional AI: Compressing Preferences into Principles." ICLR 2025.

**The claim the final round rests on:** if a written set of principles really captures what someone wants, a model reading only those principles should reproduce their choices. That is why the tool synthesizes a policy from everything a person wrote and picked, then freezes it and scores whether it agrees with them on cases they have not seen.

## Robinson — the probing technique

---

Robinson. "Probing in qualitative research interviews: theory and practice." Qualitative Research in Psychology 20, 3 (2023).

**The counterfactual probe** shown beside each case, which changes one detail to see whether the reasoning holds. Driscoll cites this as their interview method and we credit both. Ours is weaker than

theirs by design: it is one line fixed when the case is written, and cannot adapt to what the person just said.

## What we did not take

---

**GATE and OPEN** motivated the project but contributed no mechanism. Their elicitation is model-driven, which is the thing our design departs from.

**Injecting a published taxonomy into a generation prompt** looks borrowed from Botender and is not: it was already in the tool before we read them. Recorded here for honesty in the other direction.

**No human validation.** Botender ran a 90-participant study on its case generation and iRULER had two experts rate 96 essays. We have run neither.