

What we took, and from where

The papers behind the Scenario Elicitor, and concretely what each one gave the tool. The decision-by-decision record is in [provenance.md](#); this is only the things that materially shape what a participant sees.

What each paper gave us

PAPER	WHAT WE TOOK	WHERE YOU SEE IT
Driscoll	The eight topics a participant chooses from, their Table 3 verbatim, including the count of how many of the 24 parents raised each	The topic menu
Driscoll	The four things the rubric scores, synthesised from their coded moderation themes. Interpretation rather than extraction: no table in the paper contains these four	"How your rule scores"
Driscoll	The four dimensions a comparison varies: clarity of intent, severity, developmental cues, and where the risk originates	The "varies" tag on each moment
iRULER	The rubric's shape: named criteria, four ordinal levels each described as observable evidence, and a why-not-higher explanation on every score	The score column and its expanders
iRULER	Their 100% acceptance rate for AI-suggested revisions, taken as a warning rather than a feature. We built the revision and then removed it	There is no suggest-a-revision button
Ang	The five reflective question types every question in the tool is generated from. The type is recorded with each answer and never shown, since it names our taxonomy rather than the participant's task	Every question in the tool
PolicyPad	Two-layer cases: a short synopsis plus the actual conversation, rather than a description of a situation on its own	The case tabs
PolicyPad	The rule reaches the model as a system prompt over the product defaults, with the parent's instructions winning where they disagree	Not visible
PolicyPad	Write, test, revise, test again. Both of their testing paths: regenerate the reply, or continue the conversation	"Try it on a case"

PAPER	WHAT WE TOOK	WHERE YOU SEE IT
Botender	The quality bar every generated case must clear: drop one the person's stated wishes already settle, nothing that turns on tone alone, each understandable from its own text, and within what the agent can actually do	Case generation
Botender	Non-action counts as a legitimate answer, so the tool does not only ever elicit restriction	Baseline replies and options
ICAI	The claim the scored round rests on: if a written policy really captures what someone wants, a model reading only that policy should reproduce their choices	The final round
Robinson	The counterfactual probe beside each case, which changes one detail to see whether the reasoning holds. Ours is fixed when the case is written and cannot adapt the way a live interviewer would	"As you write, consider"

Full citations

SHORT NAME	CITATION
Driscoll	Driscoll, Chen, Shi, Vucharatavintara, Yao, Jin. "Understanding Parents' Desires in Moderating Children's Interactions with GenAI Chatbots through LLM-Generated Probes." CHI '26. 10.1145/3772318.3791622
iRULER	Bai, Cheong, Muller, Lim. "iRULER: Intelligible Rubric-Based User-Defined LLM Evaluation for Revision." CHI '26. 10.1145/3772318.3790539
Ang	Ang, Gollapalli, Ng. EACL 2023, Table 1. Reproduces the Paul and Elder question taxonomy
PolicyPad	Feng, Kuo, Chen, Cheong, Holstein, Zhang. "PolicyPad: Collaborative Prototyping of LLM Policies." CHI '26. 10.1145/3772318.3791689
Botender	Kuo, Liu, Chen, Seering, Zhang, Zhu, Holstein. "Botender: Supporting Communities in Collaboratively Designing AI Agents through Case-Based Provocations." CHI '26. 10.1145/3772318.3790500
ICAI	Findeis, Kaufmann, Hullermeier, Albanie, Mullins. "Inverse Constitutional AI: Compressing Preferences into Principles." ICLR 2025
Robinson	Robinson. "Probing in qualitative research interviews: theory and practice." Qualitative Research in Psychology 20, 3 (2023)